

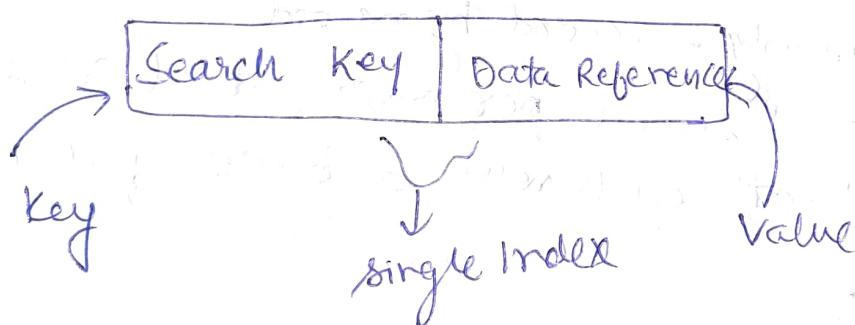
## Unit - 3

### Storage strategies

#### Indices :

- Indexing in DBMS is used to speed up data retrieval by minimizing disk scans.
- Instead of searching through all rows, the DBMS uses index structures to quickly locate data using key values.
- When an index is created, it stores sorted key values and pointers to actual data rows.
- This reduces the no. of disk accesses, improving performance especially on large datasets.

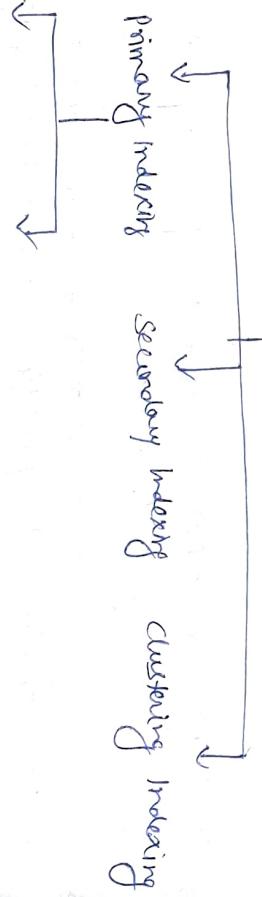
#### Structure of Index in DB



#### Attributes of Indexing :

1. Access Types - types of access, such as value based search, range based access.
2. Access Time
3. Insertion Time
4. Deletion time
5. space overhead - additional space required by the index.

## Indexing



**How Indexing makes DB faster?**

Let us takes an example we have 32 records of data:

### without indexing:

32 rows of data searching done linearly - record by record suppose it takes 1 second per row, total 32 sec, and it will treat each row a block of data.

### with indexing:

is done in sorted order based

→ in indexing 'search' is done in sorted order based on index values.

Let's imagine, 32 rows grouped into blocks, 8 rows per block → total 4 blocks

Index will store just one entry per block.

Block 1 → key 1 instead of 32 seconds,

Block 2 → key 9 it might take  $\log_2(32)$  less

Block 3 → key 17

Block 4 → key 25 it reduces no. of disk read

Ex.	age	id
21	2	
22	3	
23	1	
24	4	
25	6	

each entry in index is  
 $4B + 4B = 8B$

There will be 100 entries in the index, so total size of index will be  $8 \times 100 = 800B$

1 disk block = 600B

for 800B = 2 disk block

Query : Find all users with age = 23 with index

### Plan:

1) iterate index

    ↳ block by block

2) check age == 23 in entries

3) if yes, add the 'id' in a buffer

4) if no, discard

5) for all the relevant id in the buffer for the relevant records from the disk

    ↳ read the records from the disk

    ↳ add an off buffer

    ↳ read the records from the disk

    ↳ add an off buffer

    ↳ read the records from the disk

    ↳ add an off buffer

    ↳ read the records from the disk

    ↳ add an off buffer

    ↳ read the records from the disk

    ↳ add an off buffer

    ↳ read the records from the disk

    ↳ add an off buffer

    ↳ read the records from the disk

    ↳ add an off buffer

    ↳ read the records from the disk

    ↳ add an off buffer

    ↳ read the records from the disk

    ↳ add an off buffer

    ↳ read the records from the disk

    ↳ add an off buffer

    ↳ read the records from the disk

    ↳ add an off buffer

    ↳ read the records from the disk

    ↳ add an off buffer

## Advantages of Indexing

1. faster Queries

2. Efficient Access

3. Improved Sorting

4. consistent performance

5. Data Integrity

## Disadvantages?

1. Increased storage space

2. Increased maintenance overhead

3. slower insert/update operations

4. complexity in choosing the right index

## # B-Trees

→ specialized m-way tree designed to optimize data access especially on disk based storage.

→ In a B-tree of order m, each node can have upto m children and m-1 keys, allowing it to efficiently manage large datasets.

→ value of m is decided based on disk block and key sizes.

→ standout feature of b-tree:

1) Ability to store a significant number of keys within a single node, including large key values.

→ Height when the B-trees is completely full

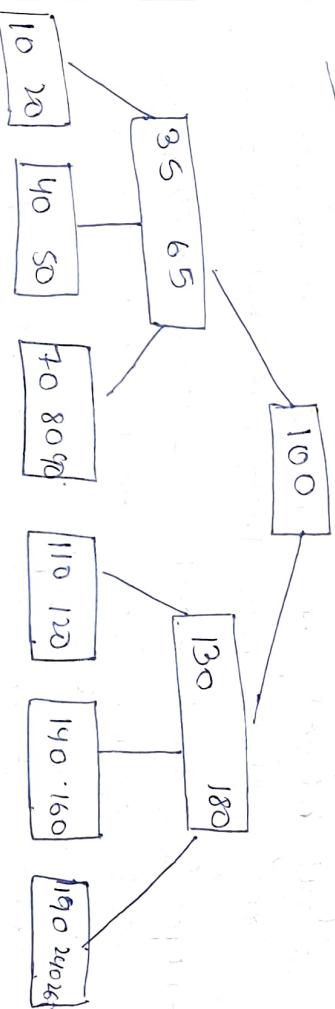
$$h_{\min} = \lceil \log(n+1) \rceil - 1$$

→ Height when the B-trees is least filled

$$h_{\max} = \lfloor \log t \left( \frac{n+1}{2} \right) \rfloor$$

## Example:

### B-Tree of order 5



### Properties of B-Tree :

1. All the leaf nodes of a B-tree are at the same level i.e. they have the same depth (height of tree)

2. The keys of each node of a B-tree, should be stored in ascending order.

3. All non-leaf nodes (except root node) should have at least  $m/2 - 1$  keys.

5. If root node is a leaf node = no children

if root node is non-leaf node = 2 children (atleast one key)

6. A non-leaf node with  $m-1$  key values should have at least one key

## Need of a B-tree :

- Improved performance over M-way trees, because it is self balanced. It makes suitable for external storage systems where fast data retrieval is crucial.
- Optimized for large datasets - it can handle millions record efficiently. Their reduced height and balanced structure make it happen.
- It has high concurrency & high - throughput
- Operations on B - Tree
  - ↳ fast & scalable in heavy load

## Operation Time Complexity

Operation	Time complexity
Search	$O(\log n)$
Insert	$O(n \log n)$
Delete	$O(n \log n)$
Traverse	$O(n)$

B tree is generated every time stored on disc

from leaf reading up  
holding in  
node reading

Update - Update, can overwrite within the same width.  $O(N) - TC$

Find one - linear scan  $O(n) - TC$

Range Queries - possible only when rows are ordered by it.  $O(m) - TC$

Delete - Create a new file without that entry/row.

$O(N) - TC$

leaf node traversal  $O(n)$   
leaf node are linked so as to enable linear traversal.

## Facts :

- If 1 B+ tree node is 4000 B big and size of document size is 40B then each node = 100 rows (max)
- size of B+ tree node  $\approx$  disk block size
- ↳ in one disk read we read 1 node  $\approx$  100 rows

# Hashing

Internal structure of Hash Table:

Hash Tables → one of the most widely used data structures

→ Hash tables are also used as building blocks for

- 1) classes and its members
- 2) variable lookup tables

→ Hash Tables are designed to provide constant time → insertions deletion lookups

String                  Integer

## Naive Implementation

Because we are already getting  $\text{int}^{\circ}$  from  $\text{key}(k)$ , what if we store the value ( $v$ ) at index  $k$ ?



we get  $O(1)$  insert, update & lookup isn't this great?

This approach works well, only when  $N$  is small  
space required for holding array when

$$N=10 \rightarrow 4 \times 10 \approx 40 \text{ B}$$

$$N=100 \rightarrow 4 \times 100 \approx 400 \text{ B}$$

$$N=1000 \rightarrow 4 \times 1000 \approx 4 \text{ KB}$$

$$N=1M \rightarrow 4 \times 1M \approx 4 \text{ MB}$$

$$\text{if } N = \text{int}^{32} \text{ range } \rightarrow 4 \times 4 \text{ billion } \approx 16 \text{ GB}$$

\* For some native types, the hash function is internally implemented.

Hash keys have a larger range say  $[0, 2^{32}]$  and hash function converts the object to this int range.

$$\text{apple} \rightarrow f \rightarrow 12762179$$

$$\text{banana} \rightarrow f \rightarrow 51962$$

$$\text{cat} \rightarrow f \rightarrow 62$$

$$\text{dog} \rightarrow f \rightarrow 1962719$$

## ~~# Application keys to Hash Keys~~

→ we cannot put anything as a key in hash table.

→ it is limited to a specific set of types

e.g. string, int, tuple etc.

→ we can also use custom types as keys if they implement the 'hash' function that splits out an integer for the object.

### Challenges:

1. finding this big chunk of memory is tough
2. lots of slot would remain empty

Adding more keys  
 if we add more keys to our hash table, the holding array to be resized.  
 $m \rightarrow 2m$

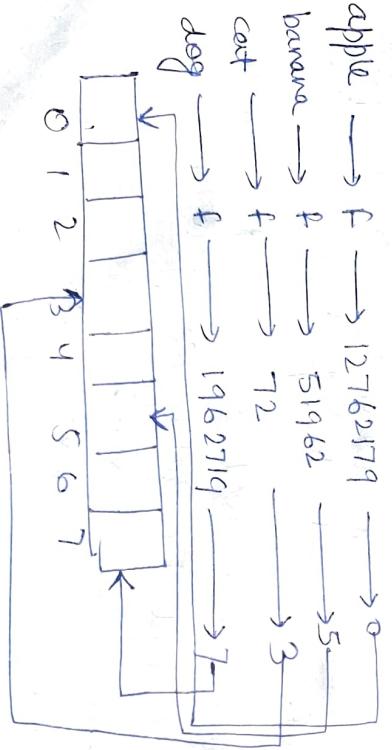
### # Mapping hash keys to smaller range

→ if we are planning to store  $K$  keys in the hash table, we have to have an array of size  $m$ , such that

$$m \in O(K)$$

This requires us to have a second step that reduces hash key from range to smaller bin range  
 $[0, N] \longrightarrow [0, m]$

Say, we are planning to store 4 keys in hash table we can have a bin array of length 8



### conflicts

$$\text{key1} \rightarrow f \rightarrow h_1$$

$$\text{key2} \rightarrow f \rightarrow h_2$$

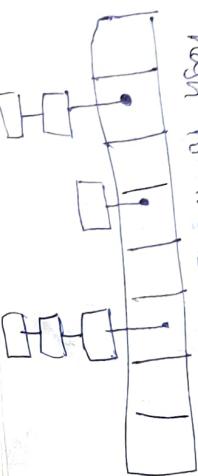
multiple keys can produce same hash key upon hashing.

so, how can we store multiple keys in the same slot?

The two classical ways of achieving this

1. Chaining
2. Open Addressing

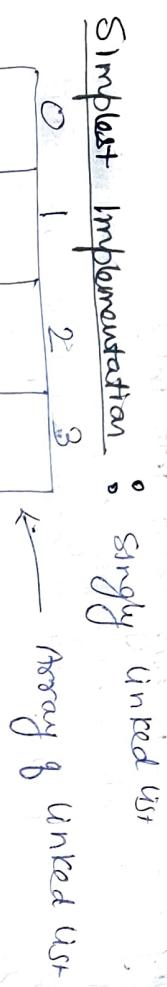
Core idea:  
 from a chain of keys that hash to the same slot



## Chaining:

We put colliding keys in a data structure that hold them well.

most common implementation  $\rightarrow$  linked list



Each slot contains the pointer to the head of the linked list. Struct slots

Struct node \* head;

$\rightarrow$  Each node of the list contains:

1. pointer to the actual key
2. pointer to the next node of the list

Struct {

    void \* key;

    Struct node \* next;

}

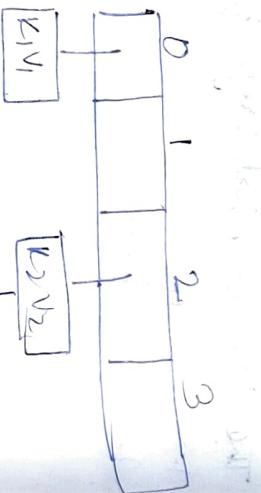
## Hash Table Operations

1. Adding a key

    put ( $k_1, v_1$ )

    put ( $k_2, v_2$ )

    put ( $k_3, v_3$ )



### 2. Delete a key

Delete  $k_2, k_1, k_3$

Delete operation is simple

1. Reach the slot in  $O(1)$
2. Iterate through the list and bind node  $\rightarrow$   $key == k_1$
3. While iterating keep track of prev node
4. Adjust the pointers
5. Delete the intended node

### 3. Lookup a key

get  $k_1, k_2, k_3$

lookups similar to delete

1. reach the node in  $O(1)$

2. linearly iterate through the list until we find node  $\rightarrow$   $key == k_1$

Given a key and a value, we can pass the key through the hash function and get index  $i$ .

1. Pass the key through the hash function and get index  $i$ .

2. Create a new linked list node with  $k, v$  at index  $i$

3. Add it to the chain present at index  $i$

## Other data structure for hashing



Instead of linkedlist, we can use a self balancing binary trees to store collected  $K, V$  pairs.

→ insertion are not O(1) but looks O(n)

## Good probing function

- The probing function should generate the permutation of numbers  $[0, m-1]$  so as to cover the entire space eventually.

## Implementing probing function

when keys collide, find a way to hunt for an available slot in the array — deterministically.

X	✓	X	X
---	---	---	---

Probing : finding the next available slot called *posting*.

Probing strategy can be defined as

$j = p(K, i)$  that spits out the new index where key  $K$  can be placed in  $i^{\text{th}}$  attempt.

$$j \in [0, m] \quad i \in [0, m]$$

$m = \text{size of hash table}$

$$\begin{array}{ccccccc} 0 & 1 & 2 & 3 & 4 & 5 & 6 & 7 \\ | & | & | & X & | & | & X & | & X \end{array}$$

So, while looking for key  $K_1$ , we first look at  $p(K_1, 0) = 5$  if we can't find,  $p(K_1, 1) = 7$ , if we can't find,  $p(K_1, 2) = 2$

$$j = p(K, m-1), \text{ if that is occupied}$$

try

Hence, first insert  $j = p(K, 0)$ , if that is occupied  
 $j = p(K, 1)$ , if that is occupied with  
 $j = p(K, 2)$ , if that is occupied we

## Hash Table Operations : Adding a key

→ Until we find a free slot, keep probing and check at the first free slot, put the key.

## Hash Table operations : lookup

- Similar to adding. Using probing function, we try to find keys in slots.
- Iterations stops when we find the key or we exhaust iterating over all the slots.

## Hash Table operations : Deleting a key

The deletion is a soft delete. We lookup the key using probing function and upon discovering, we mark the slot is deleted.

### But why soft delete?

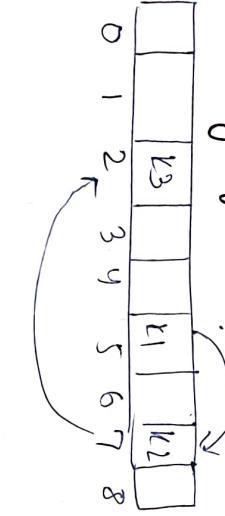
The deletion is a soft delete. We lookup the key using probing function and upon discovering, we mark the slot is deleted.

## Hash Table operation (Adding a key)

Say, we have delete key  $k_2$  and now we are looking for  $k_3$ .

DELETE  $k_2$

	$k_3$		$k_1$	$k_2$
0	1	2	3	4



$k_1, k_2, k_3$  hashed at index 5, with open addressing, say we get them placed at 5, 7 and 2.

## # Linear Probing

$$p(k, i) = (h(k) + i) \% m$$

We search linearly from the hashed index until the end of the table and then wrap from the start.

	$x$		$x$	$x$	$x$	$x$	$k_1$
0	1	2	3	4	5	6	7

## Hash Table operation (Adding a key)

- We involve the probing function to find the slot in the hash table.
- If that slot is occupied, we traverse the hash table and find first available slot.

It is like linear search from slot

If you delete  $k_2$ , then how would we ever reach  $k_3$ ?

empty slot == Stop iteration

so, we will never be able to reach  $k_3$  during lookup hence, we need to differentiate b/w free and delete.

↓  
soft deletion is the way to go.

## Hash Table Operations : Key Lookup

→ It gives a constant time performance.

### Challenges

- we invoke the probing function to get the slot  $h[k_4] = 2$
- if key present at that index, we return
- if not, we traverse to the next & try to find the key
- if we encounter the key, we return

→ If we reach the end, we start from index 0 we stop iteration as soon as we encounter an empty slot

0	1	2	3	4	5	6

$k_3$  hashes to 3

- Hash Table Operations : Deleting a key
- Delete is a soft delete that we could continue reading to the keys stored further.

Say, Del  $k_2$  then  $h(k_4) = 2$

0	1	2	3	4	5	6
		$k_1$	$k_3$	$x$	$k_4$	

Because  $k_1$  and  $k_4$  collided,  $k_3$ 's primary slot got occupied

0	1	2	3	4	5	6
		$k_1$	$k_2$	$k_3$		

Hence, good uniform hash function is essential for linear probing to be efficient.

### # Quadratic Probing

- Instead of placing the collided key in the neighbouring slot, quadratic probing adds successive value of an arbitrary quadratic polynomial.
- When we access a memory location, the page is cached on CPU and the page contains neighbouring elements.
- Subsequent accesses are served from CPU cache

Quadratic

Linear

Sample sequence:

$$h(k), h(k) + 1^2, h(k) + 2^2, h(k) + 3^2, \dots$$

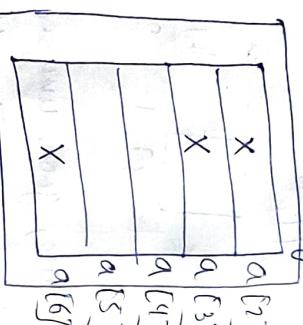
How is it better than linear probing?

- It reduces clustering and cascaded collisions as collision keys are placed further away from each other.

→ Properties of quadratic probing:

1. It reduces clustered collision by distributing it quadratically.
- \* It is not immune to it, but still it reduces it to a good extent.

2. It has a good locality of reference but not as great as linear probing.
- \* leverage CPU cache well



Unless there are large collisions on same key, at least a couple of subsequent slots will be in CPU cache.

→ Two slots are already cache on CPU.

$$h(k) = 2$$

$$h(k) + 1^2 = 3$$

$$h(k) + 2^2 = 6$$

## # Double Hashing

→ Double hashing technique uses 2 hash functions to find the slot.

→ First hash function gives primary slot and upon collision, it uses second hash offset times attempt until an empty slot is reached.

$$p(k, i) = (h_1(k) + i * h_2(k)) \bmod m$$

→ Given that we are using another hash function to offset, we are minimizing repeated collisions and effects of clustering.

→ Following no specific pattern and gives near-uniform yet random offset from the index (primary slot).

$$\text{Sequence: } h_1(k), h_1(k) + 1, h_2(k), h_1(k) + 2, h_2(k), \dots$$

→ Choosing the second hash function

1. It should never return 0
2. It should cycle through entire table (order doesn't matter)
3. Fast to compute

→ Advantages:

1. Uniform spread over collision
2. Following no specific offset pattern, purely dependent on key
3. least prone to clustering problem.  
→ offset from primary slot is uniformly distributed.