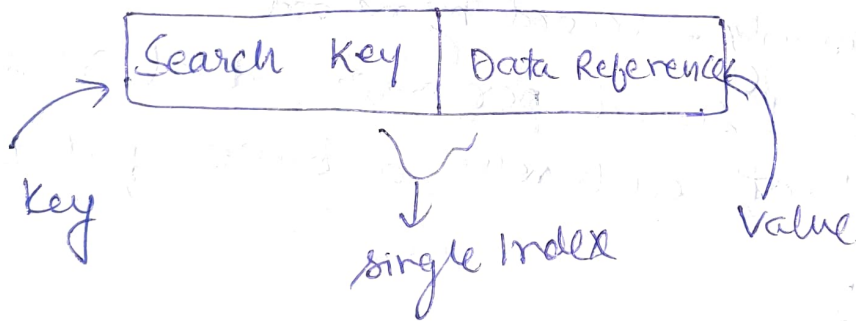
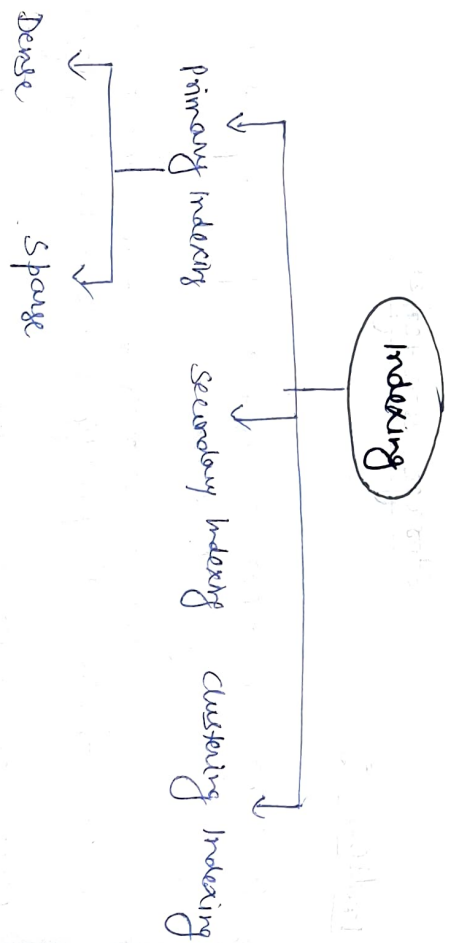


Storage strategiesIndices :

- Indexing in DBMS is used to speed up data retrieval by minimizing disk scans.
- Instead of searching through all rows, the DBMS uses index structures to quickly locate data using key values.
- When an index is created, it stores sorted key values and pointers to actual data rows.
- This reduces the no. of disk accesses, improving performance especially on large datasets.

Structure of Index in DBAttributes of Indexing :

1. Access Types — types of access, such as value based search, range based access.
2. Access Time
3. Insertion Time
4. Deletion time
5. space overhead — additional space required by the Index.



## How indexing makes ~~DB~~ faster?

Let us take an example we have 32 records of data:

### without indexing:

32 rows of data  
Searching done linearly - record by record  
Suppose it takes 1 second per row, total 32 sec,  
and it will treat each row a block of data.

### with indexing:

→ in indexing, storing is done in sorted order based on index values.

→ Let's imagine, 32 rows grouped into blocks, 8 rows per block → total 4 blocks.

Index will store just one entry per block.

Block 1 → key 1  
Block 2 → key 9  
Block 3 → key 17  
Block 4 → key 25

it reduces no. of disk read.

Ex:

age	id
21	2
22	3
22	5
23	1
23	4
24	6

1 disk block = 600B

for 800B = 2 disk block

Query: find all users with age == 23 with index

Plan:

↳ Iterate index

↳ block by block

1) check age == 23 in entries

2) if yes, add the 'id' in a buffer

3) if no, discard

4) for all the relevant id in the buffer

↳ read the records from the disk

↳ add an o/p buffer

Read the index  
not the relevant  
ids meeting  
criteria

for the relevant  
ids fetch  
the actual  
records from  
the disk.

each entry in index is  
4B + 4B = 8B

There will be 100 entries in  
the index, so total size of index  
will be  $8 \times 100 = 800B$

## Advantages of Indexing:

1. Faster Queries
2. Efficient Access
3. Improved Sorting
4. Consistent performance
5. Data Integrity

## Disadvantages:

1. Increased storage space
2. Increased maintenance overhead
3. Slower insert/update operations
4. Complexity in choosing the right index

## # B-Trees

→ specialised m-way tree designed to optimize data access, especially on disk based storage.

→ In a B-tree of order m, each node can have upto m children and m-1 keys, allowing it to efficiently manage large datasets.


→ value of m is decided based on disk block and key sizes.

→ standard feature of b-tree:

- 1) Ability to store a significant number of keys within a single node, including large key values.

## Example:

B-Tree of order 5



## Properties of B-Tree:

1. All the leaf nodes of a B-tree are at the same level i.e. they have the same depth (height of tree)
2. The keys of each node of a B-tree, should be stored in ascending order.

3. All non-leaf nodes (except root node) should have at least  $m/2 - 1$  keys.

5. If root node is a leaf node, = no children  
At least one key

if root node is non-leaf node = 2 children (at least)

6. A non-leaf node with m-1 key values should have n non null children.  
At least one key

→ Height when the B-trees is completely full

$$h_{min} = \lceil \log(n+1) \rceil - 1$$

→ Height when the B-trees is least filled

$$h_{max} = \lceil \log_t \left( \frac{n+1}{2} \right) \rceil$$



## Need of a B-tree:

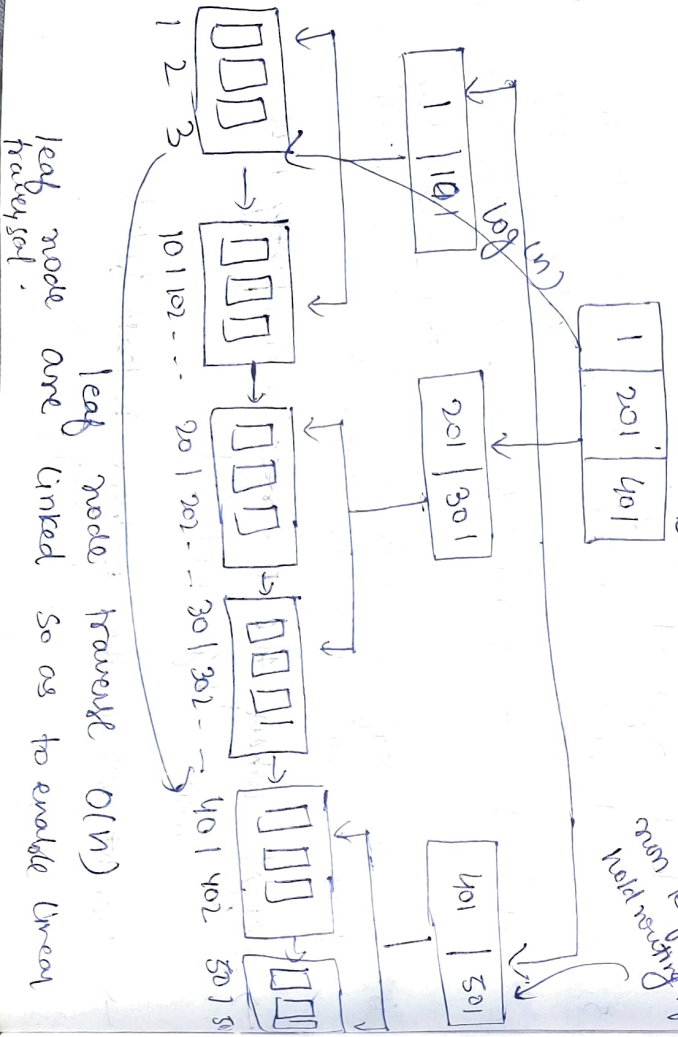
- Improved performance over M-way trees, because it is self balanced, it makes suitable for external storage systems where fast data retrieval is crucial.
  - Optimized for large datasets - it can handle millions record efficiently. Their reduced height and balanced structure make it happen.
  - It has high concurrency & high-throughput
- Operations on B-Tree

### Operation | Time complexity

Search	$O(\log n)$
Insert	$O(\log n)$
Delete	$O(\log n)$
Traverse	$O(n)$

if tree is garbled, every stored on disk

non leaf nodes holding pointers



## Facts:

- If 1 B+ tree node is 4000 B big and row/document size is 40B then each node = 100 rows (max)
- size of B+ tree node  $\approx$  disk block size
- in one disk read we read 1 node  $\approx$  100 rows

## Q: Why we we required to use B-tree?

⇒ Say our table/records are stored in one file sequentially.

Insert - cannot efficiently insert in the middle  
 $O(n)$  - TC

Update - Update, can overwrite within the same width.  $O(n)$  - TC

Find one - linear scan  $O(n)$  - TC

Range Query - Possible only when rows are ordered by it.  $O(n)$  - TC

Delete - Create a new file without that entry/rows.  $O(n)$  - TC

# Hashing

Internal structure of Hash Table:

Hash Tables  $\rightarrow$  one of the most widely used data structure

$\rightarrow$  Hash tables are also used as building block for constructing:

- 1) classes and its members
- 2) variable lookup tables

$\rightarrow$  Hash Tables are designed to provide

- Constant time  $\rightarrow$  insertions
- deletion
- lookups

Key	$\rightarrow$ Value
apple	5
banana	15
cat	18
dog	3

- Two ideas to construct hash table:
1. Application key to hash key [O(n)]  
apple  $\rightarrow$  12762179
  2. Hash key to a smaller key  
12762179  $\rightarrow$  17

## Application keys to hash keys

$\rightarrow$  we cannot put anything as a key in hash table.

$\rightarrow$  it is limited to a specific set of types  
e.g. string, int, tuple etc.

$\rightarrow$  we can also use custom types as keys if they implement the 'hash' function that splits out an integer for the object.

\* For some native types, the hash function is internally implemented.


Hash keys have a larger range say  $[0, 2^{32}]$  and a hash function converts the object to this int range.

apple  $\rightarrow$  12762179  
banana  $\rightarrow$  51962  
cat  $\rightarrow$  62  
dog  $\rightarrow$  1962719

String  $\rightarrow$  Integer

## Naive Implementation

Because we are already getting int (i) from key(k), what if we store the value (v) at index i?



This approach works well, only when N is small space required for holding array when

N=10  $\rightarrow$   $4 \times 10 \approx 40B$   
N=100  $\rightarrow$   $4 \times 100 \approx 400B$   
N=1000  $\rightarrow$   $4 \times 1000 \approx 4KB$   
N=1M  $\rightarrow$   $4 \times 1M \approx 4MB$   
if N = int32 range  $\rightarrow 4 \times 4 \text{ billion} \approx 16GB$

### Challenges:

1. Finding this big chunk of memory is tough
2. lots of slots would remain empty

### # Mapping hash keys to smaller range


→ if we are planning to store  $K$  keys in the hash table, we have to have an array of size  $m$ , such that

$$m \in O(K)$$

This requires us to have a second step that reduces hash key from range to smaller km range

$$[0, N] \rightarrow [0, m]$$

Say, we are planning to store 4 keys in hash table we can have a km [array] of length 8



### Adding more keys

if we add more keys to our hash table, the holding array to be resized.

$$m \rightarrow 2m$$

### Why are even hashing strict into int?

The first step is simplifying our problem statement for the second step, making it easier to optimize int  $\rightarrow$  int distribution.

The first step also allows us to give great abstraction enabling us to support complex data types as key.

$$\text{Object} \rightarrow \text{INT32}$$

### conflicts

$$\text{key1} \rightarrow f \rightarrow h_1$$

$$\text{key2} \rightarrow f \rightarrow h_2$$

multiple keys can produce same hash key upon hashing.


So, how can we store multiple keys in the same slot?

### The two classical ways of achieving this

1. Chaining
2. Open Addressing

#### Core idea?

from a chain of keys that hash to the same slot






## Chaining:

We put colliding keys in a data structure that hold them well.  
most common implementation  $\rightarrow$  linked list

### Simplest Implementation: Singly linked list



$\rightarrow$  Each node of the list contains:

1. pointer to the actual key
2. pointer to the next node of the list

Struct {

void \* key;


Struct node \* next;

};

### Hash Table Operations

#### 1. Adding a key

put (K<sub>1</sub>, V<sub>1</sub>)  
put (K<sub>2</sub>, V<sub>2</sub>)  
put (K<sub>3</sub>, V<sub>3</sub>)



- Given a key and a value, we
1. Pass the key through the hash function and get index 'i'.
  2. Create a new 'linked list' node with K, V
  3. Add it to the chain present at index 'i'

### Possible implementations

1. Insertion can always happen at the head  $\rightarrow$  best
2. Insertion can always happen at the tail  $\rightarrow$  fast
3. Insertion can happen as per the sort order  $\rightarrow$  linear iteration

#### 2. Delete a key

Delete K<sub>2</sub>, K<sub>1</sub>, K<sub>3</sub>

Delete operation is simple

1. Reach the slot in O(1)
2. Iterate through the list and find node  $\rightarrow$  key == K<sub>1</sub>
3. While iterating keep track of prev node
4. Adjust the pointers
5. Delete the intended node


#### 3. Lookup a key

GET K<sub>1</sub>, K<sub>2</sub>, K<sub>3</sub>

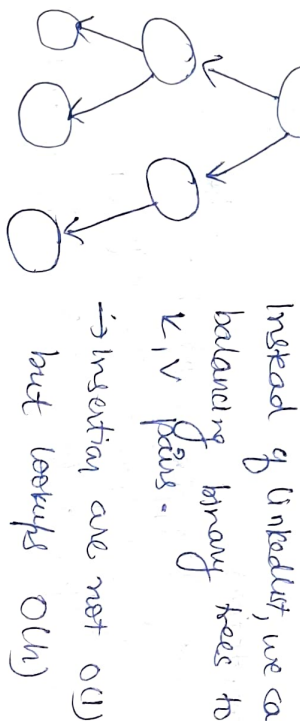
lookups similar to delete

1. reach the node in O(1)
2. linearly iterate through the list until we find node  $\rightarrow$  key == K<sub>1</sub>

## Other data structure for chaining.




Instead of linkedlist, we can use a self balancing binary trees to store collected  $k, v$  pairs.



→ insertion are not  $O(1)$  but lookups  $O(\ln)$

## # Open Addressing :

when keys collide, find a way to hunt for an available slot in the array — deterministically.



**Probing:** Finding the next available slot called probing.

Probing strategy can be defined as  $j = P(k, i)$  that spits out the next index where key  $k$  can be placed in  $m$  attempt  $i$ .

$j \in [0, m]$   $i \in [0, m]$

$m = \text{size of hash table}$

Hence, first insert

$j = P(k, 0)$ , if that is occupied we try

$j = P(k, 1)$ , if that is occupied we try

$j = P(k, 2)$ , if that is occupied we try

$j = P(k, m-1)$ , if that is occupied we try

## Good Probing function

→ The probing function should generate the permutation of numbers  $[0, m-1]$  so as to cover the entire space eventually.

## Implementing probing function

→ It is a simple mathematical or algorithmic function that deterministically tells us our next slot for a particular key.

Key attempt

$P(k, 0) \rightarrow 5$  1st attempt probing func gave 5

$P(k, 1) \rightarrow 7$  2nd attempt probing func gave 7

$P(k, 2) \rightarrow 2$  3rd attempt probing func gave 2



So, while looking for key  $k_1$ , we first look at  $P(k_1, 0) = 5$  if we can't find,  $P(k_1, 1) = 7$ , if we can't find,  $P(k_1, 2) = 2$



## Hash Table Operations: Adding a key

→ Until we find a free slot, keep probing and check at the first free slot, put the key.

## Hash Table operations: lookup


→ Similar to adding. Using probing function we try to find keys in slots.

→ Iteration steps when we find the key or we exhaust, iterating over all the slots.

## Hash Table operations: Deleting a key

The deletion is a soft delete. We loop the key using probing function and upon discovering, we mark the slot as deleted.


### But why soft delete?



Say, we hard delete key  $k_2$  and now we are looking for  $k_3$ .

DELETE  $k_2$

GET  $k_3$



If you delete  $k_2$ , then how would we ever reach  $k_3$ ?

empty slot == stop iteration

So, we will never be able to reach  $k_3$  during lookup. Hence, we need to differentiate b/w free and delete.

↓  
soft deletion is the way to go.

## Limitations of open Addressing

Max number of keys = # slots in array